

## DB2


Forum: <https://forum-db.informatik.uni-tuebingen.de/c/ss20-db2>

## Assignment 7 (16.06.2020)

Submission: Tuesday, 23.06.2020, 10:00 AM



Relevant videos: up to DB2 - Chapter 09 - Video #43.

<https://tinyurl.com/DB2-2020>1. [8 Points] B<sup>+</sup>Tree - InsertFigure 1: A B<sup>+</sup>Tree

In this assignment, let us assume that all key values inserted into the B<sup>+</sup>Tree are unique. For this exercise, further assume the index' **Insert**-operation to **not implement redistribution**.

Answer the following questions, each based on the unmodified B<sup>+</sup>Tree of Figure 1:

- Identify *four* leaf node entries ( $a, \dots, d$ ) which, when inserted one after another (**Insert**  $a, \dots, \text{Insert } d$ ), fill the leaf level pages completely.
- How many **Insert** operations are at least needed to grow the size of the tree by two levels?
- Take all entries of the leaf level sequence set (2,3,5,...) and **Insert** them, one after another, into a new B<sup>+</sup>Tree of order  $o = 1$ . Hand in sketches of the B<sup>+</sup>Tree instances after the first three **Insert** steps, together with the finally resulting B<sup>+</sup>Tree.

**Note:** You may hand in your solutions as plain text, PDF or PNG image file.

## 2. [12 Points] B<sup>+</sup>Tree - Maintenance

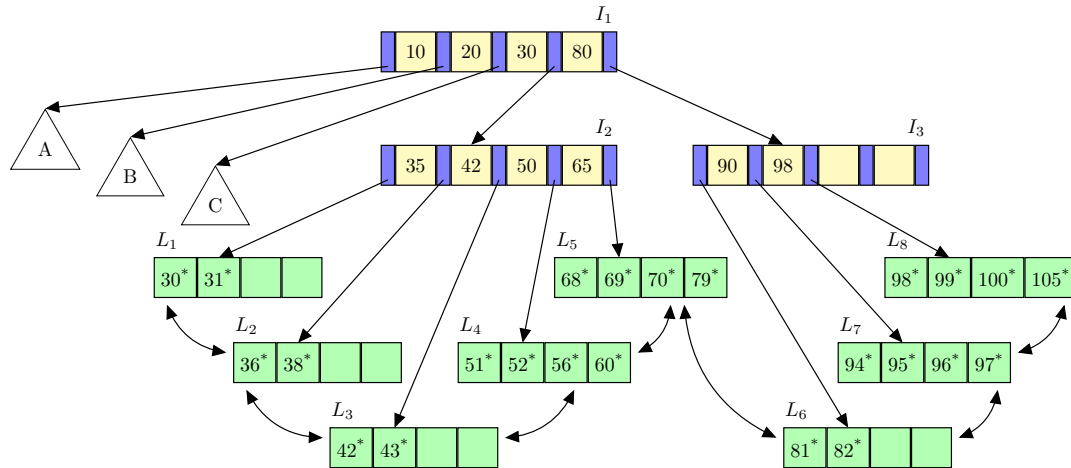


Figure 2: Another B<sup>+</sup>Tree

Again, let us assume that all key values inserted into the B<sup>+</sup>Tree are unique. However, for this exercise assume both index operations, **Insert** and **Delete**, to implement **redistribution on leaf level**. Remember that redistribution is only performed between direct siblings, i.e. directly linked leaf nodes having the same parent node.

Answer the following questions, each based on the unmodified B<sup>+</sup>Tree of Figure 2:

- Write down all nodes ( $I_j$  or  $L_k$ ) which have to be read by the RDBMS to answer the following queries:
  - “Find all records with a key value greater than 38”
  - “Delete the record with the key value 81”
- Add a record with the key value 104 to the tree. Sketch the resulting tree.
- Name all key values which, on insert, will directly cause the tree to increase in depth.
- Delete the entry with key value 36 from the tree. Sketch the resulting tree.
- The subtrees *A*, *B* and *C* have not been specified completely. Describe any characteristics which can still be inferred about these subtrees.

**Note:** You may hand in all tree sketches as plain text, PDF or PNG image file.

## 3. [10 Points] B<sup>+</sup>Tree - PostgreSQL

We provided you with a SQL file `btree.sql` which creates and populates a table

```
indexed (a INT PRIMARY KEY, b TEXT, c NUMERIC(3,2)).
```

Load the file with PostgreSQL and answer the following questions. Please hand in all SQL queries and intermediate results you used to find the final answer.

**Note:** The tasks require you to use functions `bt_metap(relname TEXT)`, `bt_page_stats(relname TEXT, blkno INT)` and `bt_page_items(relname TEXT, blkno INT)` previously mentioned in the lectures. For more information about these functions, read the documentation at

<https://www.postgresql.org/docs/current/pageinspect.html#id-1.11.7.31.6>.

- How many pages have been created for the index that PostgreSQL automatically created based on the primary key?
- Write a query to find the *root node* of the B<sup>+</sup>Tree. Which page number does it have and what is its and fan-out?
- What is the *average* fan-out of all non-leaf nodes?

- (d) Use the functions `bt_page_stats(...)` and `bt_page_items(...)` to manually traverse the B<sup>+</sup>Tree from the root to the index leaf page that holds the key entry for  $a = 150\,000$ . What are the **minimal** and **maximal key values** found on that leaf node? On which pages of relation `indexed` do the entries of that leaf node point?

**Note:** Function `bt_page_items(...)` returns a column `data` which represents the key value of an index entry. The `data` value is encoded as a hexadecimal string. For example, the value `'77 97 01 00 00 00 00 00'` is read as the hexadecimal number `0x19777` which can be converted to the decimal number 104311.

For your convenience, we provided you with a function `data_to_numeric(TEXT)` in `btree.sql`, to convert these `data` values to decimal numbers. It can be used as follows:

```
SELECT data_to_numeric('77 97 01 00 00 00 00 00');
```

```
data_to_numeric
-----
              104311
(1 row)
```