

DB2

Forum: <https://forum-db.informatik.uni-tuebingen.de/c/ss20-db2>

Assignment 9 (30.06.2020)

Submission: Tuesday, 07.07.2020, 10:00 AM



Relevant videos: up to DB2 - Chapter 11 - Video #60.

<https://tinyurl.com/DB2-2020>

1. [15 Points] Performance Impact of Indexes

Obviously, indexes have impact on the performance of data modifying operations. But, can an additional index also **decrease** the performance of a **SELECT**-query?

Consider the following table:

```
CREATE TABLE skewed (  
    sort      INTEGER NOT NULL,  
    category  INTEGER NOT NULL,  
    interesting BOOLEAN NOT NULL  
);  
  
INSERT INTO skewed  
    SELECT i AS sort, i % 1000 AS category, i > 50000 AS interesting  
    FROM generate_series(1, 1000000) i;  
  
CREATE INDEX skewed_category ON skewed (category);  
ANALYZE skewed;
```

We now want to query the first twenty interesting rows in category 42:

```
SELECT *  
FROM skewed  
WHERE interesting AND category = 42  
ORDER BY sort  
LIMIT 20;
```

- (a) Print the query plan using **EXPLAIN ANALYZE** and describe in your own words, how query evaluation proceeds and how this is supported by index **skewed_category**.

In general, it is a good idea to support top-*N* queries (**ORDER BY/LIMIT**) by an index on the sort criteria. Let us test that!

- (b) Create an additional B⁺Tree index **skewed_sort** on column **sort** of table **skewed**.
- (c) You will notice that index **skewed_sort** will decrease the performance of our query. **But why?** Again, print the query plan, explain the evaluation strategy in your own words and compare the estimated costs to the plan of **1a**. Why is it – in general – a good idea to use the index on **sort** to support this query? What is the issue that causes the performance drain in the concrete example?
- (d) Can you think of another index which employs the idea of using the sort order, but does not suffer the problem of **1c**? Find the index which fits the given query best. Measure and explain its benefit compared to **1a**.

2. [7 Points] UB-Trees: B⁺Tree on Z-order values

One possibility to index a table on multiple columns are *composite indexes*. However, these do not support both dimensions equally. Instead, the first dimension dominates the order of entries.

Consider a table `points(x INT, y INT)` representing points in a two-dimensional space:

```
-- A two dimensional space
CREATE TABLE points (x INT, y INT);


-- Populate space with points
INSERT INTO points
  (SELECT x, y
   FROM generate_series(0,99) AS x, generate_series(0,99) AS y);

-- Create primary key index
CREATE UNIQUE INDEX points_x_y ON points USING btree (x,y);
```

The *composite index* reflects spacial locality of points in a two-dimensional space only on the first dimension `x`. Locality regarding *both* dimensions can be achieved with an *expression index* that combines both dimensions in a single locality-preserving value, called *Z-order value*:

```
CREATE INDEX points_z_value_x_y ON points USING btree (z_order_value(x,y));
ANALYZE;
```

The Z-order value of a point (x,y) is calculated using the bit-interleaving function provided in `zorder.sql`.¹ This reflects locality for both input dimensions:



For both B⁺Tree indexes, `points_x_y` and `points_z_value_x_y`:

- Choose an arbitrary leaf page. Uses function `bt_page_stats(indexname, pageno)` provided by extension `pageinspect`² to identify a leaf page.
- Collect all points (x,y) referred by this leaf page. Use function `bt_page_items(indexname TEXT, pageno INT)` provided by `pageinspect` to access the *RIDs* of all items on the B⁺Tree page.
- Illustrate the (x,y) location of these points in the two-dimensional space, plotting them on a grid of size 100×100 with *Gnuplot*.

You can access a web-based version of Gnuplot at <http://gnuplot.respawned.com>.³

- Compare the two plots: describe and explain your findings **briefly**.

¹See Wikipedia for more information on Z-order values: https://en.wikipedia.org/wiki/Z-order_curve

²<https://www.postgresql.org/docs/current/pageinspect.html#id-1.11.7.31.6>

³Examples for "Data" and "Plot script" can be found in files `data.txt` and `points.plot`.

3. [8 Points] **External Merge Sort**

Given a input table of unsorted values distributed over 16 pages. Each page covers up to two elements:

pages 1-8:	38,58	23,31	36,59	21,35	19,53	13,25	56,22	54,21
pages 9-16:	34,58	19,14	48,46	47,32	60,40	58,60	55,24	50,54

Sort the table using **External Merge Sort**, with an available working memory of **B = 3 pages**.
For each *Pass* of the algorithm, write down the *output runs* written to secondary memory.

Note: *Pass 0* does **not** use *Replacement Sort* and returns sorted runs of size **B** pages.