

DB 2

05 – Row Updates

Summer 2020

Torsten Grust
Universität Tübingen, Germany

1 : Q_4 — Row Update

SQL probe Q_4 uses SQL DML statements (**INSERT**, **UPDATE**, **DELETE**) to alter the state of a table:

| | | |
|----------------------------|------------------------|----------------------------|
| INSERT INTO ternary | UPDATE ternary | DELETE FROM ternary |
| SELECT ... | SET $c = e_1$ | WHERE $a = e_2$ |
| FROM ... | WHERE $a = e_2$ | |

INSERT: evaluate query to construct new rows.

UPDATE, **DELETE**: query the table and identify the affected rows.

Modify table storage to reflect the row updates.¹

¹ We still assume that the table has *no* associated index structures.



Using **EXPLAIN** on Q_4 : **INSERT**

EXPLAIN VERBOSE

```
INSERT INTO ternary(a,b,c)
  SELECT t.a, 'Han Solo', t.c
 FROM   ternary AS t;
```



QUERY PLAN

```
Insert on public.ternary (cost=0.00..20.00 rows=1000 width=44)
-> Seq Scan on public.ternary t (cost=0.00..20.00 rows=1000 width=44)
    Output: t.a, 'Han Solo'::text, t.c
```

- **Seq Scan** scans table **ternary** to construct rows to be inserted, feeds 1000 rows into **Insert** for insertion.
- Width of inserted rows (over-)estimated to be 44 bytes = 4 (int) + 32 (text) + 8 (float) bytes.

Reading Complex **EXPLAIN** Outputs

- **EXPLAIN** uses symbol `->` and indentation to visualize larger, **tree-shaped** query evaluation plans:



- Read plans “inside out”, root ***op*₁** delivers query result.



Using **EXPLAIN** on Q_4 : **UPDATE**

EXPLAIN VERBOSE

```
UPDATE ternary AS t
SET    c = -1
WHERE  t.a = 982;
```

QUERY PLAN

```
Update on public.ternary t (cost=0.00..22.50 rows=1 width=51)
-> Seq Scan on public.ternary t (cost=0.00..22.50 rows=1 width=51)
    └─> Output: a, b, '-1'::double precision, ctid
        Filter: (t.a = 982)
```

- **Seq Scan** emits complete rows (only **c** was updated).
- Additionally feeds row ID (**ctid**) into **Update** to identify the affected row(s).



Using **EXPLAIN** on Q_4 : **DELETE**


EXPLAIN VERBOSE

```
DELETE FROM ternary AS t
WHERE t.a = 982;
```

QUERY PLAN

```
Delete on public.ternary t (cost=0.00..22.50 rows=1 width=6)
-> Seq Scan on public.ternary t (cost=0.00..22.50 rows=1 width=6)
    Output: ctid ◀
    Filter: (t.a = 982) ▶
```

- **Seq Scan** returns affected row IDs (of 6 bytes each) only.
- We turn to **Filter** (makes scan skip non-qualifying rows) later in this course.



2 : How do Row Updates Alter the Table Storage?

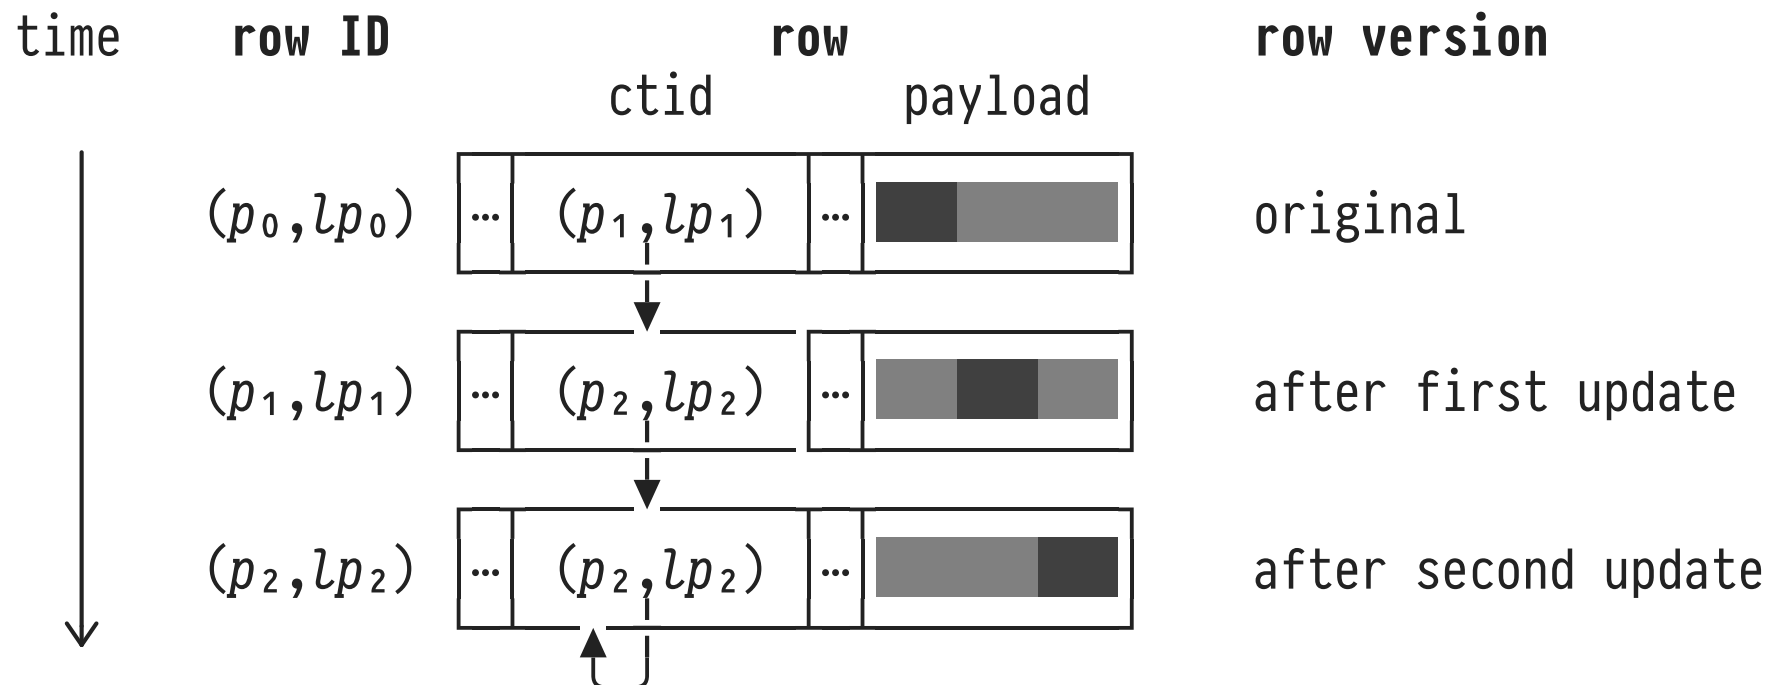
Let us take a closer look at how plan operator **Update** alters the target table's heap file pages. We find:²

- Rows are *not* updated in-place. A **new version of the row is created**—original and updated row co-exist.
- Any database user (query, application) sees exactly one version of any row at any time. Different users may see different row versions.
- A separate **VACUUM** (“garbage collection”) step collects and removes old versions that cannot be seen by any user.

² This implementation of **Update** is typical for all DBMS that implement **Multi-Version Concurrency Control (MVCC)**. We discuss MVCC later in this course.



Row Version Chains




- Original and updated versions of a row form a **chain**, linked by the rows' IDs (held in row header field **ctid**).



Row Visibility and Timestamps

1. Each row carries two **timestamps**—`xmin` and `xmax`—that mark its first and last time of existence.
2. Each query/update is executed at some timestamp `T` which defines the rows that are **visible** to the operation:



row is visible for any operation
with timestamp $t_0 \leq T < t_1$

($t_1 = \infty$: row has not been updated yet)

- DBMS uses system-wide virtual timestamps (*transaction IDs*), see PostgreSQL built-in function `txid_current()`.



Impact of Updates Beyond the Row's Page

- Updates on full pages may lead to row relocation across pages: versions then have row IDs (p_i, lp_i) , (p_{i+1}, lp_{i+1}) where $p_i \neq p_{i+1}$.
 - Traversal of longer update chains may lead to I/O-costly “page hopping.”
 - \Rightarrow Perform **VACUUM** to collect inaccessible old versions. From outside page, point to most recent row directly.
- PostgreSQL optimizes for the good-natured case where $p_i = p_{i+1}$ and indexed row fields have *not* been changed.
 - Such **heap-only tuple (HOT) updates** have page-internal impact only, no maintenance outside page required.

3 | Q_4 — Row Update

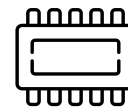


| | | |
|---|---|--|
| INSERT INTO ternary VALUES (...,...,...) | UPDATE ternary SET c = e_1 WHERE a = e_2 | DELETE FROM ternary WHERE a = e_2 |
|---|---|--|

UPDATE: affects updated column(s) only.

INSERT, **DELETE**: operate on full rows, all column BATs of table **ternary** will be affected.

- MonetDB uses user-specific **Δ tables** (“delta tables”) to represent changes. Column BATs are *not* modified immediately. Global visibility of changes is delayed.

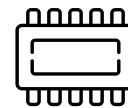


Using **EXPLAIN** on Q_4 : **DELETE**

```
sql> EXPLAIN DELETE FROM ternary WHERE a = 981;

:
ternary      :bat[:oid] := sql.tid(sql, "sys", "ternary");
a0           :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
deleted_rows :bat[:oid] := algebra.thetaselect(a0, ternary, 981:int, "==");
sql_delete   := sql.delete(sql, "sys", "ternary", deleted_rows);
:
```


- `algebra.thetaselect(b_1, b_2, v, θ)` returns the oids of those rows r in `algebra.projection(b_2, b_1)`, for which predicate `tail(r) θ v` holds.
- `sql.delete(...)` modifies the BAT of currently visible rows (obtainable via `sql.tid(...)`) for table `ternary`.
 ⚠ However, no column BAT is changed yet.



Using EXPLAIN on Q_4 : INSERT

```
sql> EXPLAIN INSERT INTO ternary(a,b,c) VALUES (1001, 'Han Solo', -2);  
  
:  
sql_append := sql.append(sql, "sys", "ternary", "a", 1001:int);  
sql_append := sql.append(sql, "sys", "ternary", "b", "Han Solo");  
sql_append := sql.append(sql, "sys", "ternary", "c", -2:dbl);  
:
```

- For `ternary(a,b,c)`, a row insert translates into three individual operations on the column BATs.
- `sql.append(...)` saves the inserted value in the Δ^i table associated with each column BAT.
⚠ The column BATs do not change yet.



Using **EXPLAIN** on Q_4 : **UPDATE**

```
sql> EXPLAIN UPDATE ternary SET c = -1 WHERE a = 982;
```

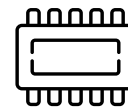
```

:
ternary      :bat[:oid] := sql.tid(sql, "sys", "ternary");
a0           :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
updated_rows :bat[:oid] := algebra.thetaselect(a0, ternary, 982:int, "==");
updated_rows_a:bat[:int] := algebra.projection(updated_rows, a0);
updated_c    :bat[:dbl] := algebra.project(updated_rows_a, -1:dbl);
sql_update   := sql.update(sql, "sys", "ternary", "c",
                           updated_rows, updated_c);
:

```

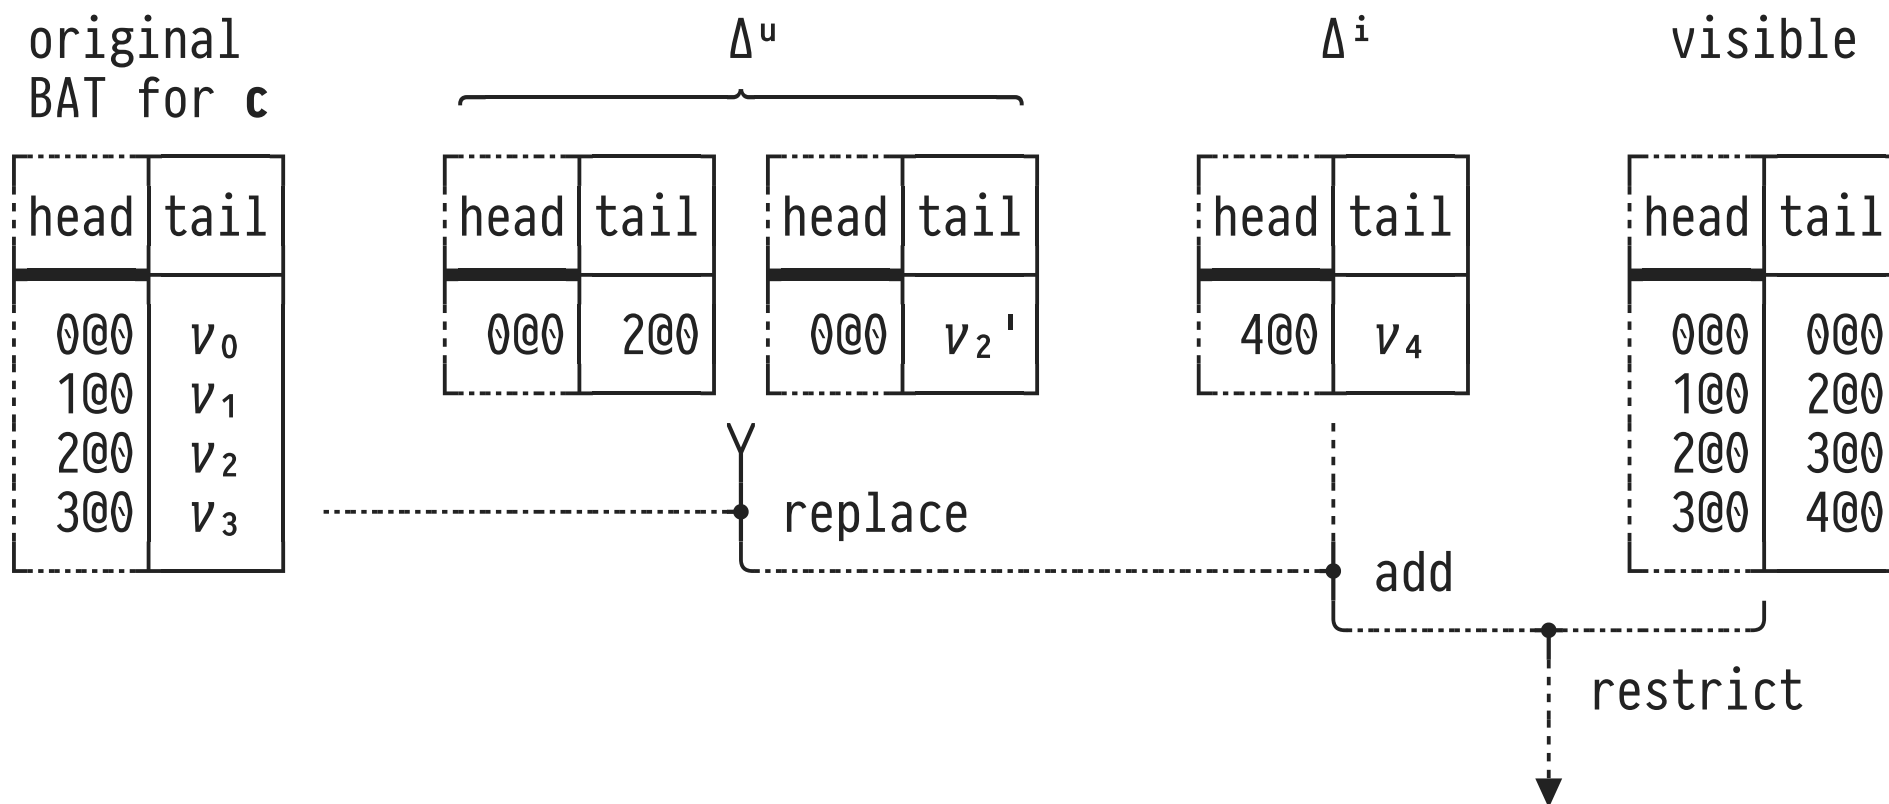
- BATs **updated_rows** and **updated_c** contain oids and **c** values of the changed rows.³
- **sql.update(...)** saves these changes in the **Δ^u table** for the BAT of column **c**. **⚠** The column BAT is not changed yet.

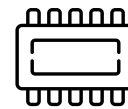
³ **algebra.project(b,v)** returns **b** with all tail values set to **v**.



Δ and Visibility Tables

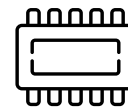
In column c , update v_2 to v_2' , insert v_4 , delete v_1 :





Applying Changes on Demand

- When a query needs to see the changes made to column `c`, apply all changes accumulated in the `c`'s Δ tables:
 1. Load—yet unmodified—column BAT for `c`.
 2. Read `c`'s Δ^u table and perform value replacements (via `bat.replace(...)`).
 3. Read `c`'s Δ^i table and perform value inserts (via `bat.append(...)`).
 4. Restrict `c` to currently visible rows (via `algebra.projection(...)`).
- Make changes permanent only once we want them to be seen globally by all users (*COMMIT* \rightarrow transaction management).



Delay Change Propagation: Disable *AUTO COMMIT*

- To experiment with Δ -based change management, disable MonetDB's default *auto commit* behavior (`mclient` option `--autocommit` or command `\a` to turn *auto commit* off):

```
DELETE FROM ternary WHERE a = 981;  
INSERT INTO ternary(a,b,c) VALUES (1001, 'Han Solo', -2);  
UPDATE ternary SET c = -1 WHERE a = 982;
```

- Without *auto commit*, changes are still pending at this point. Thus:

```
SELECT t.c      } reads  $\Delta^u$ ,  $\Delta^i$  (+ visibility table)  
FROM   ternary AS t; } to reflect changes on t.c
```